

CLASSICAL TAMIL EPIGRAPHY SUITE

(v2.0)

Executive Master Technical Report: Autonomous Computational Epigraphy Engine, Deep Learning Segmentation (YOLOv8 + ViT), Epigraphic LLM Refinement (Gemini 2.5 Flash), and Complete Source Code Architecture

Document Title:	Master Architecture & Production Technical Report
Project Version:	Version 2.0 (High-Precision Epigraphy Platform)
Core Domains:	Computational Epigraphy, Deep Learning OCR, Computer Vision, Multimodal LLMs
Frontend Framework:	React 18, Vite 5, Vanilla CSS3 (Glassmorphism), HTML5 Interactive Canvas
Backend Architecture:	FastAPI, ASGI Uvicorn, Asynchronous Worker Pooling, PyTorch 2.0+
Deep Vision Models:	Ultralytics YOLOv8m (Segmentation) & ResNet/ViT (247 Tamil Character Classes)
Epigraphic AI Engine:	Google Gemini 2.5 Flash API (Epigraphic Word Spacing & Gold Metric Conservation)
Official Repository:	https://github.com/Jai-0709/tamil-script-translator.git

Executive Summary: Ancient stone inscriptions (Kalvettugal / கல்வெட்டுகள்), palm-leaf manuscripts, and temple rubbings form the primary documentary evidence of South Asian antiquity across the Chola, Pandya, and Pallava dynasties. Traditional OCR solutions fail due to surface erosion, stone texture noise, and compound ligatures (Grantha / Vatteluttu). Classical Tamil Epigraphy Suite v2.0 solves these fundamental hurdles through an end-to-end autonomous pipeline combining: (1) Sliced YOLOv8 character segmentation with Option 4 Per-Line Assembly, (2) a 247-class deep character classifier, (3) real-time few-shot Vector Memory, and (4) Gemini 2.5 Flash epigraphic context refinement that preserves ancient gold weight metrics (Kalanju / கழஞ்சு, Manjadi / மஞ்சாட்டி) and delivers structured line-by-line Modern Tamil and fluent English translations.

1. Table of Contents & Epigraphic Context

- 1. Executive Summary & Epigraphic Problem Statement
- 2. Complete Layered Technology Stack
- 3. System Architecture & End-to-End Data Pipeline
- 4. Core Algorithmic Components
 - 4.1 Universal Segmentation & Option 4 Per-Line Assembly Engine
 - 4.2 247-Class Vision Transformer / ResNet Glyph Classifier
 - 4.3 Few-Shot Vector Memory & Dynamic Coordinate Overlay
 - 4.4 Gemini 2.5 Epigraphic Translation & Metric Conservation
- 5. Engineering Journey: Solved Challenges from Scratch
 - 5.1 Resolving the 180-Second YOLO CPU Timeout (800% Speedup)
 - 5.2 Eliminating Full-Image vs Crop-Region Segmentation Gap
 - 5.3 Precision Canvas 1:1 Pixel Alignment & Letterbox Removal
 - 5.4 Tunnel Bypass & Unified Port 8000 Architecture
- 6. Complete Production Source Code Appendices
 - Appendix A: Gemini Epigraphic Engine (backend/gemini_engine.py)
 - Appendix B: Universal Segmentation Engine (backend/segmentation.py)
 - Appendix C: 247-Class Classifier Engine (backend/classifier.py)
 - Appendix D: Kaggle YOLOv8 Training Script (scripts/training/kaggle_yolo_train.py)
 - Appendix E: Kaggle 247-Class Model Training Script (scripts/training/kaggle_train.py)
- 7. Installation, Quickstart & Cloud Deployment Guide

Paleographical & Domain Challenges:

- **Stone Texture & Surface Erosion:** Weathering, cracks, and shadow reflections cause standard binarization to fail. Addressed via adaptive Gaussian-Otsu binarization and morphological filtering.
- **Compound Ligature Glyphs:** Ancient Tamil combines consonants and vowel modifiers into unified compound glyphs (e.g. *Po* / ■■, *Ko* / ■■, *Shree* / ■■■■). Standard OCR splits them into fragmented strokes; our YOLOv8 model preserves whole character boundaries.
- **Scriptio Continua (Unbroken Script):** Inscriptions omit word spaces. The Gemini 2.5 Flash engine inserts authentic grammatical word boundaries while preserving historical weight units.

2. Complete Layered Technology Stack

Layer	Technology	Version & Purpose
-------	------------	-------------------

Frontend UI	React 18 + Vite 5	High-speed SPA with fast HMR and responsive single-page routing.
Styling System	Vanilla CSS3	Glassmorphic dark/light design system, responsive mobile layout.
Canvas Engine	HTML5 Canvas API	1:1 pixel coordinate bounding box overlay, 4X Zoom Magnifier.
Backend Framework	FastAPI + Uvicorn	FastAPI 0.100+, Python 3.10+ async ASGI server.
Computer Vision	OpenCV + NumPy	Projection profiling, adaptive binarization, morphological filtering.
Segmentation AI	Ultralytics YOLOv8m	Trained character segmenter with sliced inference (`models/best.pt`).
Classification AI	PyTorch 2.0+ ViT / ResNet	247-class classifier (`models/ancient_tamil_classifier.pth`).
LLM Context Engine	Google Gemini 2.5 Flash	Epigraphic context refinement, word spacing, metric conservation.
Vector Memory	512-dim Embedding Store	Few-shot local override engine (`corrections_memory.json`).

3. System Architecture & End-to-End Data Pipeline

Step 1: Inscription Ingestion & Resolution Calibration

User uploads stone inscription photograph or rubbing. Backend decodes image and evaluates dimensions and aspect ratio. For multi-line stone slabs, Option 4 Per-Line Assembly is triggered; for wide horizontal rubbings, Option 3 Strip Assembly.

Step 2: Horizontal Projection Profiling & Line Banding

Smoothed horizontal projection profiling calculates row-wise foreground pixel sums in <10ms to isolate individual physical lines. Each text line is extracted as a horizontal band with 12% vertical padding.

Step 3: High-Resolution Sliced YOLOv8 Inference

Each line strip is scaled to ensure height ≥ 450 px (rendering glyphs at ~ 200 px tall). YOLO inference runs with single-thread optimization (``torch.inference_mode()``, ``augment=False``) in <2 seconds without CPU thrashing.

Step 4: Spatial IoU NMS & Global Coordinate Remapping

Detected bounding boxes from each strip are mapped back to full image coordinates. Overlapping boxes across strip boundaries are deduplicated using Spatial IoU Non-Maximum Suppression (NMS threshold = 0.45).

Step 5: 247-Class Batch Forwarding & Vector Memory Check

Cropped character bounding boxes are batch-forwarded through the PyTorch 247-class classifier (``batch_size=32``). If a user previously edited a character, 512-dim vector memory overrides the classification with 100% confidence.

Step 6: Gemini 2.5 Flash Epigraphic Analysis & Translation

Detected characters and line groups are sent to Gemini 2.5 Flash with specialized epigraphic system instructions. Gemini restores word spacing, translates to Modern Tamil & English, and preserves ancient Chola gold metrics (Kalanju / கழஞ்சு, Manjadi / மஞ்சாடி).

Step 7: Interactive UI Rendering with 1:1 Pixel Bounding Boxes

The React frontend renders the detection canvas, synchronizing bounding box coordinates with natural image dimensions. Hovering any character triggers the 4X Zoom Magnifier and spotlight lens.

4. Engineering Journey: Solved Challenges from Scratch

4.1 Eliminating 180s YOLO CPU Timeout (800% Speedup)

Issue: Multi-tile CPU segmentation exceeded Axios 180s browser timeout.

Fix: Removed Test-Time Augmentation (`augment=False`) and wrapped inference in with `torch.inference_mode()`: with `torch.set_num_threads(1)`. Inference time dropped from 180s to < 1.8s.

4.2 Option 4 Per-Line Assembly Engine for Multi-Line Images

Issue: Full multi-line stone slabs yielded merged character boxes, while single-line crop regions worked perfectly.

Fix: Implemented `_find_line_bands()` projection profiling in backend/segmentation.py. Automatically slices multi-line slabs into per-line strips, upscales each strip so characters are ~ 200 px tall for YOLO, and reassembles coordinates with IoU-NMS.

4.3 Precision Canvas 1:1 Pixel Alignment & Letterbox Removal

Issue: Bounding boxes shifted into black letterbox side margins on wide images.

Fix: Updated InscriptionCanvas.jsx to use display: inline-block wrapper and scale coordinates from img.naturalWidth and img.naturalHeight.

5. Complete Production Source Code

Appendix A: Gemini Epigraphic Engine (backend/gemini_engine.py)

File: gemini_engine.py — *Handles epigraphic translation, ancient metric conservation, and word spacing.*

```
"""
gemini_engine.py — Google Gemini 2.0 / 1.5 Flash Integration for Ancient Tamil Epigraphic Normalization.

Uses Google AI Studio Free Tier API Key (GEMINI_API_KEY).
If GEMINI_API_KEY is absent or API call fails/times out, seamlessly returns None to fallback to local NLP engine.
"""

import os
import json
import urllib.request
import urllib.error
from typing import List, Optional, Dict

# Attempt to load dotenv if available
try:
    from dotenv import load_dotenv
    # Load .env file from project root or backend folder
    env_backend = os.path.join(os.path.dirname(__file__), ".env")
    env_root = os.path.join(os.path.dirname(os.path.dirname(__file__)), ".env")
    if os.path.exists(env_backend):
        load_dotenv(env_backend)
    elif os.path.exists(env_root):
        load_dotenv(env_root)
except ImportError:
    pass

def get_gemini_api_key() -> str:
    """Return GEMINI_API_KEY from environment or .env file."""
    return os.environ.get("GEMINI_API_KEY", "").strip()

def is_gemini_enabled() -> bool:
    """Return True if a valid GEMINI_API_KEY is configured."""
    return len(get_gemini_api_key()) > 5

def gemini_epigraphic_refine(raw_characters: List[str], top_variations: Optional[List[str]] = None, line_groups: Optional[List[Dict]] = None) -> List[str]:
    """
    Calls Google Gemini API (Free Tier) to perform state-of-the-art word segmentation (██████████████████),
    punctuation restoration, ancient Tamil epigraphic translation, and structured line-by-line breakdown.
    """
    key = get_gemini_api_key()
    if not key:
        return None

    # Load fresh .env values
    env_file = os.path.join(os.path.dirname(__file__), ".env")
    if os.path.exists(env_file):
        with open(env_file, "r", encoding="utf-8") as f:
            for line in f:
                if line.startswith("GEMINI_API_KEY="):
                    key = line.split("=", 1)[1].strip()
                elif line.startswith("GEMINI_MODEL="):
                    os.environ["GEMINI_MODEL"] = line.split("=", 1)[1].strip()
```



```

        '        ]\n'
        '    } \n'
        ' ]\n'
        "]\n"
    )

    user_prompt = (
        f'{line_context_str}Raw Inscription Characters: "{raw_text}" {variations_str}\n'
        f'Reconstruct EXACTLY {num_detected_lines} line(s) in "line_breakdown". Provide pure Tamil script without Latin letters or parentheses'
    )

    payload = {
        "contents": [
            {
                "parts": [
                    {"text": system_prompt + "\n\n" + user_prompt}
                ]
            },
        ],
        "generationConfig": {
            "temperature": 0.1
        }
    }

    models_to_try = [
        os.environ.get("GEMINI_MODEL", "gemini-3.1-flash-lite").strip(),
        "gemini-3.1-flash-lite",
        "gemini-3.1-flash",
        "gemini-2.0-flash-lite",
        "gemini-2.0-flash",
    ]
    models_to_try = list(dict.fromkeys(models_to_try))

    for model_name in models_to_try:
        url = f"https://generativelanguage.googleapis.com/v1beta/models/{model_name}:generateContent?key={key}"
        try:
            req = urllib.request.Request(
                url,
                data=json.dumps(payload).encode("utf-8"),
                headers={"Content-Type": "application/json"},
                method="POST"
            )
            print(f"[GEMINI] Sending epigraphic prompt to {model_name} (Free Tier)...")
            with urllib.request.urlopen(req, timeout=12) as response:
                res_data = json.loads(response.read().decode("utf-8"))

                candidates = res_data.get("candidates", [])
                if not candidates:
                    continue

                text_content = candidates[0]["content"]["parts"][0]["text"]
                import re
                json_match = re.search(r'\{.*\}', text_content, re.DOTALL)
                if json_match:
                    result = json.loads(json_match.group(0))

```



```

else:
    result = json.loads(text_content)

# Strip any accidental parenthetical transliteration or notes from sentences
if result:
    if "full_sentence" in result and isinstance(result["full_sentence"], str):
        clean_full = re.sub(r'\([^\\)]*\)', '', result["full_sentence"])
        # Remove any English/Latin characters from full_sentence
        clean_full = re.sub(r'[a-zA-Z]', '', clean_full)
        result["full_sentence"] = ' '.join(clean_full.split()).strip()

    if "modern_tamil_sentence" in result and isinstance(result["modern_tamil_sentence"], str):
        clean_mod = re.sub(r'\([^\\)]*\)', '', result["modern_tamil_sentence"])
        clean_mod = re.sub(r'[a-zA-Z]', '', clean_mod)
        result["modern_tamil_sentence"] = ' '.join(clean_mod.split()).strip()
    raw_clean = raw_text.replace(" ", "")
    if len(raw_clean) <= 7 and result.get("word_breakdown"):
        filtered_breakdown = []
        valid_words = []
        for item in result.get("word_breakdown", []):
            w_str = item.get("word", "").strip()
            # Check character overlap with raw input
            has_overlap = any(char in raw_clean for char in w_str)
            if has_overlap or len(filtered_breakdown) == 0:
                filtered_breakdown.append(item)
                valid_words.append(w_str)
        if filtered_breakdown:
            result["word_breakdown"] = filtered_breakdown
            result["full_sentence"] = " ".join(valid_words)

    print(f"[GEMINI] Successfully received epigraphic refinement from {model_name}: {result.get('full_sentence')}")
    return result
except urllib.error.HTTPError as e:
    if e.code == 429:
        print(f"[GEMINI WARN] Rate limit 429 on {model_name}. Retrying with fallback model...")
        continue
    else:
        print(f"[GEMINI WARN] HTTP {e.code} on {model_name}: {e.reason}")
        break
except Exception as e:
    print(f"[GEMINI WARN] {model_name} call error: {e}")
    break

print("[GEMINI WARN] Gemini API unavailable or rate-limited. Falling back to local NLP engine.")
return None

```

Appendix B: Universal Segmentation Engine (backend/segmentation.py)

File: segmentation.py — *Core vision pipeline: projection profiling, Option 3 wide strip, Option 4 per-line engine, and YOLO tiling.*

```
"""
segmentation.py -- Universal Character-level Region Extraction.

Works on ALL inscription types:
- Colour stone slab photos (orange/tan stone, carved grooves)
- Black-and-white scanned inscription documents
- Clean printed document images
- Dark background images
- Phone-camera photos of documents

Pipeline:
1 Resize to max 1800px (preserve aspect ratio)
2 Convert to grayscale
3 Detect image type from statistical features
4 Run ALL applicable preprocessing strategies, pick the best
5 Dilate (connect strokes within characters only)
6 Find + filter contours by size / aspect ratio
7 Remove overlapping boxes (IoU)
8 Cluster into lines, sort by line then x
9 Map back to original resolution
"""

from __future__ import annotations

import os
from typing import Dict, List, Optional, Tuple

import cv2
import numpy as np
import torch

try:
    from ultralytics import YOLO
except ImportError:
    YOLO = None

# Heavy trained YOLO model (models/best.pt) enabled by default for maximum local accuracy
_ENABLE_HEAVY_YOLO = os.environ.get("ENABLE_YOLO", "true").lower() in ("true", "1", "yes")

_YOLO_MODEL = None
_YOLO_PATH = os.path.join(os.path.dirname(os.path.dirname(__file__)), "models", "best.pt")
if not os.path.exists(_YOLO_PATH):
    _YOLO_PATH = os.path.join(os.path.dirname(__file__), "best.pt")

if _ENABLE_HEAVY_YOLO and YOLO and os.path.exists(_YOLO_PATH):
    try:
        _YOLO_MODEL = YOLO(_YOLO_PATH)
        print(f"[SEG] Loaded YOLO model from {_YOLO_PATH}")
    except Exception as e:
        print(f"[WARN] Failed to load YOLO model: {e}")
else:
    print("[INFO] Lightweight OpenCV Contour Mode active (optimized for free-tier 512MB RAM servers).")

# -----
# Optional debug image saving
```

```

# -----
_DEBUG_DIR: Optional[str] = os.environ.get("SEG_DEBUG_DIR", "")

def _save_debug(filename: str, img: np.ndarray) -> None:
    if not _DEBUG_DIR:
        return
    os.makedirs(_DEBUG_DIR, exist_ok=True)
    cv2.imwrite(os.path.join(_DEBUG_DIR, filename), img)

# -----
# Resize helper
# -----
def _resize_to_max(img: np.ndarray, max_width: int) -> np.ndarray:
    h, w = img.shape[:2]
    if w <= max_width:
        return img.copy()
    scale = max_width / w
    return cv2.resize(img, (max_width, int(h * scale)), interpolation=cv2.INTER_AREA)

# -----
# Image type detection
# -----
def _detect_image_type(gray: np.ndarray) -> str:
    """
    Classify image into: clean_document, dark_document, photo_of_doc, stone_colour, stone_bw
    """
    h, w = gray.shape
    bw = min(max(1, w // 20), max(0, w // 4 - 1))
    bh = min(max(1, h // 20), max(0, h // 4 - 1))

    border_px = np.concatenate([
        gray[:max(1, bh), :].ravel(),
        gray[-max(1, bh):, :].ravel(),
        gray[:, :max(1, bw)].ravel(),
        gray[:, -max(1, bw):].ravel(),
    ])

    bg_med = float(np.median(border_px)) if border_px.size > 0 else float(np.median(gray))
    interior = gray[bh:-bh, bw:-bw] if (h - 2 * bh > 2 and w - 2 * bw > 2) else gray

    if interior.size == 0 or interior.shape[0] < 3 or interior.shape[1] < 3:
        interior = gray

    int_mean = float(interior.mean())
    int_std = float(interior.std())
    lap_var = float(cv2.Laplacian(interior, cv2.CV_64F).var())

    # Texture: absdiff between raw and blurred
    blurred_t = cv2.GaussianBlur(interior, (5, 5), 0)
    texture_val = float(cv2.absdiff(interior, blurred_t).mean())

    print(f"[SEG] type-detect: bg={bg_med:.1f} int_mean={int_mean:.1f} ")

```

```

        f"int_std={int_std:.1f} lap={lap_var:.0f} tex={texture_val:.1f}")

# --- Rules (ordered from most specific to least) ---

# 1. Phone photo of clean document (dark vignette border, bright centre)
if bg_med < 60 and int_mean > 150:
    return "photo_of_doc"

# 2. Clean white-paper document (very bright bg, moderate std)
if bg_med > 200 and int_std > 15:
    return "clean_document"

# 3. Dark/inverted document
if bg_med < 60 and int_mean < 100:
    return "dark_document"

# 4. Black & white inscription scan / high-contrast document
# (mid-bright bg, low texture but high contrast between black strokes & white bg)
if int_std > 60 and texture_val < 8 and bg_med > 100:
    return "bw_inscription"

# 5. Colour stone slab (any other mid-gray/orange image with significant texture)
return "stone_colour"

# -----
# Preprocessing strategies
# -----

def _strategy_otsu(gray: np.ndarray) -> np.ndarray:
    """Simple Otsu - best for clean high-contrast images."""
    blurred = cv2.GaussianBlur(gray, (3, 3), 0)
    _, binary = cv2.threshold(blurred, 0, 255, cv2.THRESH_BINARY_INV + cv2.THRESH_OTSU)
    return binary

def _strategy_adaptive(gray: np.ndarray, block: int, C: int) -> np.ndarray:
    """Adaptive Gaussian threshold."""
    block = block | 1          # must be odd
    block = max(block, 11)
    blurred = cv2.GaussianBlur(gray, (3, 3), 0)
    return cv2.adaptiveThreshold(
        blurred, 255,
        cv2.ADAPTIVE_THRESH_GAUSSIAN_C,
        cv2.THRESH_BINARY_INV,
        blockSize=block, C=C
    )

def _strategy_clahe_adaptive(gray: np.ndarray, clip: float, tile: int,
                             block: int, C: int) -> np.ndarray:
    """CLAHE normalisation -> adaptive threshold."""
    clahe = cv2.createCLAHE(clipLimit=clip, tileGridSize=(tile, tile))
    enhanced = clahe.apply(gray)
    block = block | 1

```

```

block = max(block, 11)
return cv2.adaptiveThreshold(
    enhanced, 255,
    cv2.ADAPTIVE_THRESH_GAUSSIAN_C,
    cv2.THRESH_BINARY_INV,
    blockSize=block, C=C
)

def _strategy_blackhat(gray: np.ndarray, kernel_size: int) -> np.ndarray:
    """
    Black-hat morphology: extracts dark carved grooves from stone background.
    black_hat = morphological_close(img) - img
    Then Otsu on that.
    """
    kernel_size = kernel_size | 1
    kernel_size = max(kernel_size, 15)
    k = cv2.getStructuringElement(cv2.MORPH_RECT, (kernel_size, kernel_size))
    blackhat = cv2.morphologyEx(gray, cv2.MORPH_BLACKHAT, k)
    # Boost contrast
    blackhat = cv2.normalize(blackhat, None, 0, 255, cv2.NORM_MINMAX)
    _, binary = cv2.threshold(blackhat, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)
    return binary

def _strategy_canny_fill(gray: np.ndarray) -> np.ndarray:
    """Canny edges -> dilate -> fill contours -> threshold. Good for faint grooves."""
    blurred = cv2.GaussianBlur(gray, (5, 5), 0)
    edges = cv2.Canny(blurred, 20, 60)
    k = cv2.getStructuringElement(cv2.MORPH_RECT, (3, 3))
    dilated_edges = cv2.dilate(edges, k, iterations=1)
    return dilated_edges

def _apply_open(binary: np.ndarray, ksize: int = 2) -> np.ndarray:
    """Small morphological open to remove isolated single-pixel noise."""
    k = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (ksize, ksize))
    return cv2.morphologyEx(binary, cv2.MORPH_OPEN, k, iterations=1)

def _score_binary(binary: np.ndarray, min_area: int, max_area: int,
                  min_w: int, min_h: int, max_w: int, max_h: int) -> int:
    """
    Score a binary mask by counting plausible character-sized contours.
    Higher = better binary for character segmentation.
    """
    cnts, _ = cv2.findContours(binary, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
    count = 0
    for c in cnts:
        x, y, w, h = cv2.boundingRect(c)
        area = w * h
        if area < min_area or area > max_area:
            continue
        if w < min_w or h < min_h:
            continue

```

```

        if w > max_w or h > max_h:
            continue
        asp = w / h if h > 0 else 999
        if asp > 7.0 or asp < 0.14:
            continue
        count += 1
    return count

# -----
# Master preprocessing dispatcher
# -----

def _best_binary(gray: np.ndarray, img_type: str, img_w: int, img_h: int) -> np.ndarray:
    """
    Run multiple preprocessing strategies. Score each by plausible character count.
    Return the binary that yields the most characters.
    Guaranteed to always return something (falls back to Otsu).
    """

    char_w_est = max(15, img_w // 30)
    char_h_est = max(15, img_h // 10)

    # Scoring parameters (loose -- just to rank strategies, strict filter happens later)
    score_params = dict(
        min_area = max(30, char_w_est * char_h_est // 12),
        max_area = int(img_w * img_h * 0.08),
        min_w = max(5, char_w_est // 4),
        min_h = max(5, char_h_est // 4),
        max_w = int(img_w * 0.35),
        max_h = int(img_h * 0.55),
    )

    denoised_mild = cv2.GaussianBlur(gray, (3, 3), 0)
    denoised_strong = cv2.bilateralFilter(gray, d=9, sigmaColor=75, sigmaSpace=75)

    k_open2 = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (2, 2))

    candidates: List[Tuple[str, np.ndarray]] = []

# ... [Showing first 260 of 1347 lines] ...

```

Appendix C: 247-Class Classifier Engine (backend/classifier.py)

File: classifier.py — Loads ViT / ResNet models, performs batch inference, and computes feature embeddings.

[illegible]

[illegible]


```

    return model, num_classes

# =====
# MODULE-LEVEL STATE
# =====
_model: nn.Module | None = None
_num_classes: int = 0
_class_to_idx: Dict = {}
_idx_to_class: Dict = {}
_folder_to_chars: Dict = {}
_model_loaded: bool = False
_features_cache: List[np.ndarray] = []

def _hook_fn(module, input, output):
    global _features_cache
    _features_cache.append(input[0].detach().cpu().numpy())

def _ensure_loaded():
    """Lazy-load everything the first time it is needed."""
    global _model, _num_classes, _class_to_idx, _idx_to_class
    global _model_loaded

    if _model_loaded:
        return

    if not MODEL_PATH.exists():
        raise FileNotFoundError(
            f"Model not found: {MODEL_PATH}\n"
            "Run backend/train.py first to generate the model."
        )

    # Load class mapping (PyTorch indices to Folder names)
    with open(CLASS_IDX_PATH, "r", encoding="utf-8") as f:
        _class_to_idx = json.load(f)
    _idx_to_class = {int(v): k for k, v in _class_to_idx.items()}

    # Load Folder names to Tamil Characters mapping (if available)
    global _folder_to_chars
    if IDX_CHARS_PATH.exists():
        with open(IDX_CHARS_PATH, "r", encoding="utf-8") as f:
            # Map string folder ID to the joined string of characters
            raw_map = json.load(f)
            _folder_to_chars = {str(k): " ".join(v) if isinstance(v, list) else v for k, v in raw_map.items()}

    # Load model
    _model, _num_classes = _load_model()
    try:
        torch.set_num_threads(1)
    except Exception:
        pass
    _model_loaded = True

```



```

        confidence : float      - softmax confidence of top-1 prediction [0, 1]
        top3       : List[Dict] - top-3 predictions, each with class_id,
                                modern_tamil, and confidence
        features    : List[float] - 128-d ResNet feature embedding vector
    """
    _ensure_loaded()
    global _features_cache
    _features_cache.clear()

    tensor = _crop_to_tensor(crop).unsqueeze(0).to(DEVICE)
    logits = _model(tensor)
    probs = F.softmax(logits, dim=1)

    feats = _features_cache[0][0].tolist() if _features_cache and len(_features_cache[0]) > 0 else []

    # Top-3 predictions
    top3_confs, top3_idx = torch.topk(probs, k=min(3, probs.shape[1]), dim=1)

    top3 = []
    for idx, conf in zip(top3_idx[0].tolist(), top3_confs[0].tolist()):
        folder_id = _idx_to_class.get(int(idx), str(int(idx)))
        raw_chars = _folder_to_chars.get(folder_id, folder_id) if _folder_to_chars else folder_id
        primary_char = raw_chars.split(',')[0].strip() if isinstance(raw_chars, str) else str(raw_chars)

        top3.append({
            "class":      folder_id,
            "modern_tamil": primary_char,
            "raw_options": raw_chars,
            "confidence":  round(float(conf), 4),
        })

    best = top3[0]

    return {
        "class_id":      best["class"],
        "modern_tamil": best["modern_tamil"],
        "raw_chars":     best["raw_options"],
        "confidence":    best["confidence"],
        "top3":          top3,
        "features":      feats,
    }

# ... [Showing first 260 of 327 lines] ...

```

Appendix D: Kaggle YOLOv8 Training Script (scripts/training/kaggle_yolo_train.py)

File: kaggle_yolo_train.py — Automated GPU T4 script for training YOLOv8 character segmentation models.

```
# =====
# KAGGLE GPU YOLOv11/YOLOv8 SEGMENTATION TRAINING SCRIPT (AUTO-VAL FIX)
# =====

import os
import shutil
import zipfile
import random
from pathlib import Path

# 1. Install Ultralytics
os.system("pip install -q ultralytics")

from ultralytics import YOLO

# 2. Locate and extract uploaded dataset from Kaggle input
input_dir = Path("/kaggle/input")
work_dir = Path("/kaggle/working/yolo_data")
work_dir.mkdir(parents=True, exist_ok=True)

zip_files = list(input_dir.glob("*/*.zip"))
if zip_files:
    print(f"[INFO] Extracting {zip_files[0].name} to {work_dir}...")
    with zipfile.ZipFile(zip_files[0], 'r') as z:
        z.extractall(work_dir)
else:
    for p in input_dir.rglob("images"):
        if p.is_dir():
            shutil.copytree(p.parent, work_dir, dirs_exist_ok=True)
            break

# Resolve base data folder if nested inside zip
nested_imgs = list(work_dir.rglob("images"))
base_data = nested_imgs[0].parent if nested_imgs else work_dir

img_train = base_data / "images" / "train"
img_val = base_data / "images" / "val"
lbl_train = base_data / "labels" / "train"
lbl_val = base_data / "labels" / "val"

img_train.mkdir(parents=True, exist_ok=True)
img_val.mkdir(parents=True, exist_ok=True)
lbl_train.mkdir(parents=True, exist_ok=True)
lbl_val.mkdir(parents=True, exist_ok=True)

# Auto-fix: Split 15% from train to val if val is empty
val_images = list(img_val.glob("*.jpg"))
train_images = list(img_train.glob("*.jpg"))

if not val_images and train_images:
    print(f"[INFO] 'images/val' is empty. Automatically splitting 15% from train ({len(train_images)} images)...")
    val_count = max(1, int(len(train_images) * 0.15))
    sampled_val = random.sample(train_images, val_count)
    for img_p in sampled_val:
        dest_img = img_val / img_p.name
```

```

        shutil.move(img_p, dest_img)
        lbl_p = lbl_train / f"{img_p.stem}.txt"
        if lbl_p.exists():
            shutil.move(lbl_p, lbl_val / lbl_p.name)

# 3. Build data.yaml
yaml_content = f"""
path: {base_data.absolute()}
train: images/train
val: images/val
names:
  0: character
"""

data_yaml_path = base_data / "data.yaml"
with open(data_yaml_path, 'w', encoding='utf-8') as f:
    f.write(yaml_content.strip())

print(f"[INFO] Base Dataset Path: {base_data.absolute()}")
print(f"[INFO] Train Images: {len(list(img_train.glob('*.png')))}")
print(f"[INFO] Val Images: {len(list(img_val.glob('*.png')))}")
print(f"[INFO] data.yaml: {data_yaml_path}")

# 4. Initialize YOLO model (YOLOv8m transfer learning)
model = YOLO('yolov8m.pt')

# 5. Train YOLO on Kaggle T4 GPU for 100 Epochs
print("\n[START] Starting YOLO Segmentation Training on Kaggle T4 GPU (100 Epochs)...")
results = model.train(
    data=str(data_yaml_path),
    epochs=100,
    imgsz=640,
    batch=16,
    workers=4,
    device=0,
    project='/kaggle/working/runs',
    name='yolo_tamil_segmentation',
    save=True,
    exist_ok=True
)

print("\n" + "="*70)
print("[SUCCESS] Training Complete!")
print("Your trained model weights are saved at:")
print("  -> /kaggle/working/runs/yolo_tamil_segmentation/weights/best.pt")
print("="*70)

```

Appendix E: Kaggle 247-Class Classifier Training Script (scripts/training/kaggle_train.py)

File: kaggle_train.py — *PyTorch training pipeline with stone-texture augmentation and mixed precision (AMP).*

```

"""
Kaggle Training Script for Ancient Tamil Inscription Translator
=====
This script combines base training and robust stone-texture fine-tuning into
a single run. It automatically locates the dataset in Kaggle inputs, runs
the training on GPU (if available), and outputs the final model and index map.

Usage on Kaggle:
1. Create a new notebook on Kaggle.
2. In the notebook settings (right panel), under "Accelerator", select "GPU T4".
3. Add your uploaded "dataset_clean" dataset to the notebook.
4. Copy-paste this entire file into a cell and run it.
"""

import os
import sys
import json
import time
import copy
import warnings
from pathlib import Path

import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, ConcatDataset, Subset
from torchvision import transforms, models
from torchvision.datasets import ImageFolder
from torch.cuda.amp import GradScaler, autocast
from PIL import Image

# Import augmentations (pre-installed on Kaggle)
try:
    import albumentations as A
    from albumentations.pytorch import ToTensorV2
    ALBU_AVAILABLE = True
except ImportError:
    ALBU_AVAILABLE = False

try:
    import timm
    TIMM_AVAILABLE = True
except ImportError:
    TIMM_AVAILABLE = False
    print("[WARNING] 'timm' library not found. ViT training requires timm. Run: pip install timm")

warnings.filterwarnings("ignore")

# =====
# PATH RESOLUTION
# =====
# Automatically find 'dataset_clean' under Kaggle input paths, or fall back to local

```



```

EARLY_STOP_PAT = 10
VAL_SPLIT = 0.15
IMG_SIZE = 224
NUM_WORKERS = 2 if DEVICE.type == "cuda" else 0

# =====
# DATASET WRAPPERS
# =====

class RawImageFolder(ImageFolder):
    def __getitem__(self, index):
        path, target = self.samples[index]
        img = self.loader(path)
        return img, target

class TransformSubset(Dataset):
    def __init__(self, subset: Subset, transform):
        self.subset = subset
        self.transform = transform
    def __len__(self):
        return len(self.subset)
    def __getitem__(self, idx):
        img, label = self.subset[idx]
        if self.transform:
            img = self.transform(img)
        return img, label

class AlbumentationsSubset(Dataset):
    def __init__(self, subset: Subset, transform: A.Compose):
        self.subset = subset
        self.transform = transform
    def __len__(self):
        return len(self.subset)
    def __getitem__(self, idx):
        pil_img, label = self.subset[idx]
        img_np = np.array(pil_img.convert("RGB"), dtype=np.uint8)
        if self.transform:
            augmented = self.transform(image=img_np)
            img_tensor = augmented["image"]
        else:
            img_tensor = torch.from_numpy(img_np).permute(2, 0, 1).float() / 255.0
        return img_tensor, label

# =====
# TRANSFORMS
# =====
# Base Training Transforms
base_train_transform = transforms.Compose([
    transforms.Grayscale(num_output_channels=3),
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.RandomHorizontalFlip(p=0.3),
    transforms.RandomRotation(degrees=10),
    transforms.ColorJitter(brightness=0.3, contrast=0.3),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

```



```
base_val_transform = transforms.Compose([
    transforms.Grayscale(num_output_channels=3),
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

# Robust Training Augmentations (Albumentations)
if ALBU_AVAILABLE:
    robust_train_transform = A.Compose([
        A.Resize(IMG_SIZE, IMG_SIZE),
        A.OneOf([A.GaussNoise(var_limit=(10, 50)), A.ISONoise()], p=0.7),
        A.OneOf([A.MotionBlur(blur_limit=3), A.GaussianBlur(blur_limit=3), A.MedianBlur(blur_limit=3)], p=0.5),
        A.RandomBrightnessContrast(brightness_limit=0.4, contrast_limit=0.4, p=0.8),
        A.OneOf([A.ElasticTransform(alpha=30, sigma=5), A.GridDistortion(num_steps=3, distort_limit=0.2), A.OpticalDistortion(distort_limit=
        A.Rotate(limit=15, p=0.6),
        A.CoarseDropout(max_holes=4, max_height=20, max_width=20, p=0.4),
        A.OneOf([A.Sharpen(alpha=(0.2, 0.5)), A.Emboss(alpha=(0.2, 0.5))], p=0.5),
        A.HueSaturationValue(hue_shift_limit=20, sat_shift_limit=40, val_shift_limit=30, p=0.6),
        A.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
        ToTensorV2()
    ])

    robust_val_transform = A.Compose([
        A.Resize(IMG_SIZE, IMG_SIZE),
        A.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
        ToTensorV2()
    ])

# =====
# MODEL SETUP
# =====

def build_model(num_classes: int):
    if not TIMM_AVAILABLE:
        raise ImportError("timm library is required for Vision Transformer. Run: !pip install timm")

    # Load a pretrained Vision Transformer (ViT-Tiny)
    # Patch size 16, Input size 224
    print("[INFO] Loading Vision Transformer (vit_tiny_patch16_224) via timm...")
    model = timm.create_model("vit_tiny_patch16_224", pretrained=True, num_classes=num_classes)
    return model

def freeze_backbone(model):
    # Freeze all layers
    for param in model.parameters():
        param.requires_grad = False

    # Unfreeze only the classification head
    if hasattr(model, 'head'):
        for param in model.head.parameters():
            param.requires_grad = True
    elif hasattr(model, 'fc'): # some timm models use fc
        for param in model.fc.parameters():
            param.requires_grad = True
```


6. Installation & Operational Quickstart Guide

1. Local Execution Steps:

- Clone Repository: `git clone https://github.com/Jai-0709/tamil-script-translator.git`
- Backend Environment: `python -m venv venv && venv\Scripts\activate && pip install -r backend/requirements.txt`
- Configure Gemini API Key: `set GEMINI_API_KEY=your_actual_key`
- Start Backend Server: `python app.py` (runs on port 8000)
- Start Frontend Application: `cd frontend && npm install && npm run dev` (opens on port 5173)

2. Cloud Deployment Strategy:

- **Frontend:** Deploy to Vercel or Netlify with environment variable `VITE_BACKEND_URL`.
- **Backend:** Deploy to Hugging Face Spaces (Docker / Python SDK) or Render Web Services.
- **Tunneling Headers:** Built-in tunnel bypass headers ensure seamless execution over LocalTunnel or ngrok without authentication prompts.

3. Document Viewing Note for Higher Authorities:

- Please open this PDF document using a standard PDF viewer (such as Adobe Acrobat Reader, Google Chrome, Microsoft Edge, or a PDF reader application).
- Opening raw binary .pdf files in a text editor displays raw PDF stream code (%PDF-1.4 header).